

## Sheet 05

# PS Parallel Programming

Patrick Wintner

April 8, 2025

## 1 Missing Flush Directives

Effects of missing flush directives are observed.

### 1.1 Source Code

```
1  #include <omp.h>
2  #include <stdio.h>
3
4  int main() {
5
6      int data;
7      int flag = 0;
8
9      #pragma omp parallel num_threads(2)
10     {
11
12         if (omp_get_thread_num() == 0) {
13
14             data = 42;
15
16             flag = 1;
17
18         } else if (omp_get_thread_num() == 1) {
19
20             int flag_val = 0;
```

```

21
22     while (flag_val < 1) {
23
24         flag_val = flag;
25
26     }
27
28     printf("flag=%d data=%d\n", flag, data);
29
30 }
31
32 }
33
34 return 0;
35 }

```

The program spawns two threads. Thread 0 does some work (setting the value of the variable data) before setting a flag. Thread 1 should print the values of the flag and the variable after the other thread has finished his work.

## 1.2 Experiment Method

The experiment was done on the LCC3 cluster by calling

```
salloc -exclusive -tasks-per-node=1 -cpus-per-task=1 srun -pty bash .
```

Followed by calling

```
./main.sh
```

manually. The following scripts are involved in the experiment.

### 1.2.1 Main Script

```

1  #!/bin/bash
2  # Usage: ./main.sh
3  make
4  for i in {0..999}

```

```
5 do
6     ./ex1
7 done
8 make clean
```

### 1.3 Experiment Results

The program neither terminates nor prints any output. This is probable the case because thread 1 does not fetch the updated value of flag from shared memory. and is thus not able to leave the loop.

### 1.4 Discussion

The program does indeed require several flush directives, see the code below. Those are placed either after a variable (that is read in another thread) is written to or before a variable (that is written in another thread) is read. It does not require any atomic directives, because there is no variable to which is written in both threads.

```
1  #include <omp.h>
2  #include <stdio.h>
3
4  int main() {
5
6      int data;
7      int flag = 0;
8
9      #pragma omp parallel num_threads(2)
10     {
11
12         if (omp_get_thread_num() == 0) {
13
14             data = 42;
15             #pragma omp flush(data)
16
17             flag = 1;
18             #pragma omp flush(flag)
19
20         } else if(omp_get_thread_num() == 1) {
21
22             int flag_val = 0;
```

```

23
24     while (flag_val < 1) {
25
26         #pragma omp flush(flag)
27         flag_val = flag;
28
29     }
30
31     #pragma omp flush(data)
32     printf("flag=%d data=%d\n", flag, data);
33
34 }
35
36 }
37
38 return 0;
39 }

```

## 2 Parallelising Code Snippets

### 2.1

#### 2.1.1 Original

```

a[0] = 0;
#pragma omp parallel for
for (i=1; i<N; i++) {
    a[i] = 2.0*i*(i-1);
    b[i] = a[i] - a[i-1];
}

```

#### 2.1.2 Fixed

```

a[0] = 0;
#pragma omp parallel for ordered
for (i=1; i<N; i++) {
    a[i] = 2.0*i*(i-1);
    #pragma omp ordered {
        b[i] = a[i] - a[i-1];
    }
}

```

```
    }  
}
```

The ordered construct is necessary to guarantee that the instructions are executed strictly in order.

## 2.2

### 2.2.1 Original

```
a[0] = 0;  
#pragma omp parallel  
{  
    #pragma omp for nowait  
    for (i=1; i<N; i++) {  
        a[i] = 3.0*i*(i+1);  
    }  
    #pragma omp for  
    for (i=1; i<N; i++) {  
        b[i] = a[i] - a[i-1];  
    }  
}
```

### 2.2.2 Fixed

```
a[0] = 0;  
#pragma omp parallel  
{  
    #pragma omp for  
    for (i=1; i<N; i++) {  
        a[i] = 3.0*i*(i+1);  
    }  
    #pragma omp for  
    for (i=1; i<N; i++) {  
        b[i] = a[i] - a[i-1];  
    }  
}
```

If threads do not wait after the first loop, it is possible that some elements of b are set with values of some elements of a that are not yet initialized.

## 2.3

### 2.3.1 Original

```
#pragma omp parallel for default(none)
for (i=0; i<N; i++) {
    x = sqrt(b[i]) - 1;
    a[i] = x*x + 2*x + 1;
}
```

### 2.3.2 Fixed

```
#pragma omp parallel for default(none) private(i, x) shared(a, b) ordered
for (i=0; i<N; i++) {
    #pragma ordered {
        x = sqrt(b[i]) - 1;
    }
    a[i] = x*x + 2*x + 1;
}
```

Setting default to none requires the programmer to specify the storage attribute for each variable explicitly.

## 2.4

### 2.4.1 Original

```
f = 2;
#pragma omp parallel for private(f,x)
for (i=0; i<N; i++) {
    x = f * b[i];
    a[i] = x - 7;
}
a[0] = x;
```

### 2.4.2 Fixed

```
f = 2;
#pragma omp parallel for firstprivate(f) lastprivate(x)
for (i=0; i<N; i++) {
    x = f * b[i];
    a[i] = x - 7;
}
a[0] = x;
```

The value of private variables is undefined on entry and exit of parallel regions.

## 2.5

### 2.5.1 Original

```
sum = 0;
#pragma omp parallel for
for (i=0; i<N; i++) {
    sum = sum + b[i];
}
```

### 2.5.2 Fixed

```
sum = 0;
#pragma omp parallel for reduction(+: sum)
for (i=0; i<N; i++) {
    sum = sum + b[i];
}
```

The reduction clause generates a local copy of sum and initializes it with 0. Updates occur on local copies (this would become a problem if sum were shared), which are afterwards combined to a single value.

## 2.6

### 2.6.1 Original

```
#pragma omp parallel
#pragma omp for
for (i=0; i<N; i++) {
    #pragma omp for
    for (j=0; j<N; j++) {
        a[i][j] = b[i][j];
    }
}
```

### 2.6.2 Fixed

```
#pragma omp parallel for private(j)
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        a[i][j] = b[i][j];
    }
}
```

Avoids nesting of parallel regions. Only the loop variable of the parallel for loop is by default private, therefore it is necessary to specify j as private. Another solution is to allow nesting of parallel regions by using the runtime routine `omp_set_nested()`.