

Sheet 06

PS Parallel Programming

Patrick Wintner

April 27, 2025

1 Approximation of Pi

The execution times of different implementations of the Monte Carlo method are measured.

1.1 Source Code

1.1.1 Serial Implementation

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8  int main() {
9      long n = 700000000;
10     long i, count = 0;
11     double x, y, pi;
12     double startTime, endTime;
13
14     startTime = omp_get_wtime();
15
16     srand((unsigned) time(NULL));
```

```

17     for (i = 0; i < n; i++) {
18         x = (double) rand() / RAND_MAX;
19         y = (double) rand() / RAND_MAX;
20
21         if (x * x + y * y <= 1) count++;
22     }
23
24     endTime = omp_get_wtime();
25
26     pi = 4.0 * count / n;
27     if(pi<0.99*M_PI || 1.01*M_PI<pi) {
28         fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
29         return 1;
30     }
31     printf("%.2.4f\n", endTime-startTime);
32     return 0;
33 }

```

1.1.2 Parallel Implementation using a Critical Section

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8
9  int main() {
10     long n = 700000000;
11     long i, count = 0;
12     double x, y, pi;
13     double startTime, endTime;
14
15     startTime = omp_get_wtime();
16
17     srand((unsigned) time(NULL));
18     #pragma omp parallel for private(x, y)
19     for (i = 0; i < n; i++) {
20         x = (double) rand() / RAND_MAX;
21         y = (double) rand() / RAND_MAX;
22

```

```

23         #pragma omp critical
24         {
25             if (x * x + y * y <= 1) count++;
26         }
27     }
28
29     endTime = omp_get_wtime();
30
31     pi = 4.0 * count / n;
32
33     if(pi<0.99*M_PI|| 1.01*M_PI<pi) {
34         fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
35         return 1;
36     }
37     printf("%.2.4f\n", endTime-startTime);
38     return 0;
39 }

```

1.1.3 Parallel Implementation using an Atomic Statement

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8  int main() {
9      long n = 700000000;
10     long i, count = 0;
11     double x, y, pi;
12     double startTime, endTime;
13
14     startTime = omp_get_wtime();
15
16     srand((unsigned) time(NULL));
17     #pragma omp parallel for private(x,y)
18     for (i = 0; i < n; i++) {
19         x = (double) rand() / RAND_MAX;
20         y = (double) rand() / RAND_MAX;
21
22         if (x * x + y * y <= 1) {

```

```

23         #pragma omp atomic
24         count++;
25     }
26 }
27
28 endTime = omp_get_wtime();
29
30 pi = 4.0 * count / n;
31 if(pi<0.99*M_PI|| 1.01*M_PI<pi) {
32     fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
33     return 1;
34 }
35 printf("%.24f\n", endTime-startTime);
36 return 0;
37 }

```

1.1.4 Parallel Implementation using a reduction clause

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8  int main() {
9      long n = 700000000;
10     long i, count = 0;
11     double x, y, pi;
12     double startTime, endTime;
13
14     startTime = omp_get_wtime();
15
16     srand((unsigned) time(NULL));
17     #pragma omp parallel for reduction(+: count) private(x,y)
18     for (i = 0; i < n; i++) {
19         x = (double) rand() / RAND_MAX;
20         y = (double) rand() / RAND_MAX;
21
22         if (x * x + y * y <= 1) count++;
23     }
24

```

```

25     endTime = omp_get_wtime();
26
27     pi = 4.0 * count / n;
28     if(pi<0.99*M_PI|| 1.01*M_PI<pi) {
29         fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
30         return 1;
31     }
32     printf("%.2.4f\n", endTime-startTime);
33     return 0;
34 }

```

1.2 Measurement Method

The measurement was done on the LCC3 cluster by calling `./job.sh serial 3`, `./job.sh critical 3`, `./job.sh atomic 3` and `./job.sh reduction 3`.

The following scripts are involved in the experiment.

1.2.1 SLURM Script

```


1  #!/bin/bash
2  # usage: sbatch [slurm_options] <executable> <number_of_measurements>
3
4  # Execute job in the partition "lua" unless you have special requirements.
5  #SBATCH --partition=lua
6  # Name your job to be able to identify it later
7  #SBATCH --job-name csba4017
8  # Redirect output stream to this file
9  #SBATCH --output=output.log
10 # Maximum number of tasks (=processes) to start in total
11 #SBATCH --ntasks=1
12 # Maximum number of tasks (=processes) to start per node
13 #SBATCH --ntasks-per-node=1
14 # Enforce exclusive node allocation, do not share with other jobs
15 #SBATCH --exclusive
16
17 ./main.sh $1 $2

```

1.2.2 Main Script

```
1  #!/bin/bash
2  # Usage: ./main.sh <executable> <number_of_measurements>
3
4  results=$1".dat"
5  echo "x y ey" > $results
6  make
7  for i in {1,4,8,12}
8  do
9      measurements=$i_"$1".log"
10     export OMP_NUM_THREADS=$i
11     for j in $(seq 1 $2)
12     do
13         ./$1 >> $measurements
14     done
15     ./toTable $measurements $results $2 $i
16     rm $measurements
17 done
18 make clean
```

1.3 Experiment Results



1.4 Discussion