

## Sheet 06

# PS Parallel Programming

Patrick Wintner

April 28, 2025

## 1 Approximation of Pi

The execution times of different implementations of the Monte Carlo method are measured.

### 1.1 Source Code

#### 1.1.1 Serial Implementation

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8  int main() {
9      long n = 700000000;
10     long i, count = 0;
11     double x, y, pi;
12     double startTime, endTime;
13
14     startTime = omp_get_wtime();
15
16     srand((unsigned) time(NULL));
```

```

17     for (i = 0; i < n; i++) {
18         x = (double) rand() / RAND_MAX;
19         y = (double) rand() / RAND_MAX;
20
21         if (x * x + y * y <= 1) count++;
22     }
23
24     endTime = omp_get_wtime();
25
26     pi = 4.0 * count / n;
27     if(pi<0.99*M_PI || 1.01*M_PI<pi) {
28         fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
29         return 1;
30     }
31     printf("%.2.4f\n", endTime-startTime);
32     return 0;
33 }

```

### 1.1.2 Parallel Implementation using a Critical Section

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8
9  int main() {
10     long n = 700000000;
11     long i, count = 0;
12     double x, y, pi;
13     double startTime, endTime;
14
15     startTime = omp_get_wtime();
16
17     #pragma omp parallel private(x, y)
18     {
19         unsigned seed = (unsigned) time(NULL)+13*omp_get_thread_num();
20         #pragma omp for schedule (static)
21         for (i = 0; i < n; i++) {
22             x = (double) rand_r(&seed) / RAND_MAX;

```

```

23         y = (double) rand_r(&seed) / RAND_MAX;
24
25         if (x * x + y * y <= 1){
26             #pragma omp critical
27             {
28                 count++;
29             }
30         }
31     }
32 }
33
34 endTime = omp_get_wtime();
35
36 pi = 4.0 * count / n;
37
38 if(pi<0.99*M_PI|| 1.01*M_PI<pi) {
39     fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
40     return 1;
41 }
42 printf("%.24f\n", endTime-startTime);
43 return 0;
44 }

```

### 1.1.3 Parallel Implementation using an Atomic Statement

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8  int main() {
9      long n = 700000000;
10     long i, count = 0;
11     double x, y, pi;
12     double startTime, endTime;
13
14     startTime = omp_get_wtime();
15
16     #pragma omp parallel private(x,y)
17     {

```

```

18     unsigned seed = (unsigned) time(NULL) + 13*omp_get_thread_num();
19     #pragma omp for schedule (static)
20     for (i = 0; i < n; i++) {
21         x = (double) rand_r(&seed) / RAND_MAX;
22         y = (double) rand_r(&seed) / RAND_MAX;
23
24         if (x * x + y * y <= 1) {
25             #pragma omp atomic
26             count++;
27         }
28     }
29 }
30
31 endTime = omp_get_wtime();
32
33 pi = 4.0 * count / n;
34 if(pi<0.99*M_PI|| 1.01*M_PI<pi) {
35     fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
36     return 1;
37 }
38 printf("%.2.4f\n", endTime-startTime);
39 return 0;
40 }

```

#### 1.1.4 Parallel Implementation using a Reduction Clause

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <time.h>
4  #include <omp.h>
5  #define _USE_MATH_DEFINES
6  #include <math.h>
7
8  int main() {
9      long n = 700000000;
10     long i, count = 0;
11     double x, y, pi;
12     double startTime, endTime;
13
14     startTime = omp_get_wtime();
15
16     #pragma omp parallel private(x,y)

```

```

17     {
18         unsigned seed = (unsigned) time(NULL)+13*omp_get_thread_num();
19         #pragma omp for reduction(+: count) schedule (static)
20         for (i = 0; i < n; i++) {
21             x = (double) rand_r(&seed) / RAND_MAX;
22             y = (double) rand_r(&seed) / RAND_MAX;
23
24             if (x * x + y * y <= 1) count++;
25         }
26     }
27
28     endTime = omp_get_wtime();
29
30     pi = 4.0 * count / n;
31     if(pi<0.99*M_PI|| 1.01*M_PI<pi) {
32         fprintf(stderr, "Error: estimated value deviates significantly: %f\n", pi);
33         return 1;
34     }
35     printf("%.2.4f\n", endTime-startTime);
36     return 0;
37 }

```

## 1.2 Measurement Method

The measurement was done on the LCC3 cluster by calling `sbatch job.sh serial 3`, `sbatch job.sh critical 3`, `sbatch job.sh atomic 3` and `sbatch job.sh reduction 3`.

The following scripts are involved in the experiment.

### 1.2.1 SLURM Script

```

1  #!/bin/bash
2  # usage: sbatch [slurm_options] <executable> <number_of_measurements>
3
4  # Execute job in the partition "lua" unless you have special requirements.
5  #SBATCH --partition=lua
6  # Name your job to be able to identify it later
7  #SBATCH --job-name csba4017
8  # Redirect output stream to this file
9  #SBATCH --output=output.log

```

```

10 # Maximum number of tasks (=processes) to start in total
11 #SBATCH --ntasks=1
12 # Maximum number of tasks (=processes) to start per node
13 #SBATCH --ntasks-per-node=1
14 # Enforce exclusive node allocation, do not share with other jobs
15 #SBATCH --exclusive
16
17 ./benchmark.sh $1 $2

```

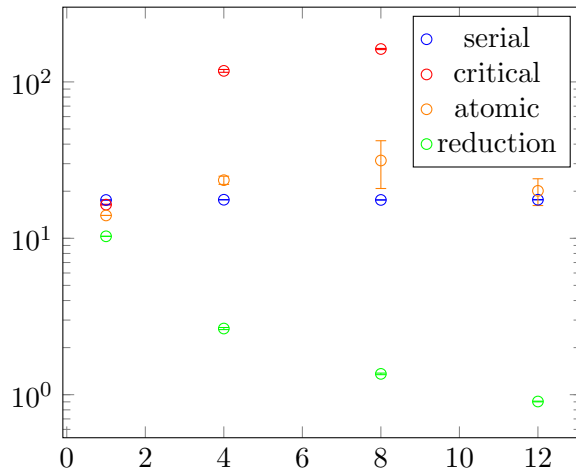
### 1.2.2 Benchmark Script

```

1  #!/bin/bash
2  # Usage: ./benchmark.sh <executable> <number_of_measurements>
3
4  results=$1".dat"
5  echo "x y ey" > $results # create header
6  make
7  for i in {1,4,8,12} # number of threads
8  do
9      measurements=$i_"$1".log"
10     export OMP_NUM_THREADS=$i
11     for j in $(seq 1 $2) # repeat measurement $2 times
12     do
13         ./ $1 >> $measurements # store measurement results in <executable>.log
14     done
15     ./toTable $measurements $results $2 $i #store table in <executable>.dat
16     rm $measurements
17 done
18 make clean

```

### 1.3 Experiment Results



### 1.4 Discussion

Unsurprisingly, there is no performance impact when using different numbers of threads for the serial implementation. The performance of the parallel implementation with an atomic statement is roughly equal to the performance of the parallel implementation and slightly deteriorates with an increasing number of threads. Considering that only the variable assignments to `x` and `y` (calling `rand_r`) runs in parallel, while incrementing the counter is done serial, this is not really surprising. The performance of the parallel implementation with a critical section deteriorates significantly when increasing the number of threads. Only the implementation using a reduction clause gains performance by increasing the number of threads. This meets the expectations, because each thread has its own count, thus not slowing each other when incrementing the counter.

## 2 Mandelbrot

The execution time of an parallel implementation using openOMP of the mandelbrot set is measured.

## 2.1 Source Code

```
1  #include <stdint.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <time.h>
5  #include <omp.h>
6
7  // Include that allows to print result as an image
8  #define STB_IMAGE_WRITE_IMPLEMENTATION
9  #include "stb_image_write.h"
10
11 // Default size of image
12 #define X 1280
13 #define Y 720
14 #define MAX_ITER 10000
15
16 void calc_mandelbrot(uint8_t image[Y][X]) {
17     size_t i, j;
18     #pragma omp parallel for schedule(guided) private(j)
19     for(i=0; i<Y; ++i) {
20         for(j=0; j<X; ++j) {
21             double x=0;
22             double y=0;
23             double cx=(double)j/(X-1)*3.5-2.5; // scale j to [-2.5, 1]
24             double cy=(double)i/(Y-1)*2-1; // scale i to [-1, 1]
25             size_t iteration=0;
26             while(x*x+y*y<2*2 && iteration<MAX_ITER) {
27                 double x_tmp=x*x-y*y+cx;
28                 y=2*x*y+cy;
29                 x=x_tmp;
30                 iteration=iteration+1;
31             }
32             // scale iteration to [0,255]
33             char norm_iteration=iteration*255/MAX_ITER;
34             image[i][j]=norm_iteration;
35         }
36     }
37 }
38
39 int main() {
40     uint8_t image[Y][X];
41 }
```



```

42     double startTime = omp_get_wtime();
43     calc_mandelbrot(image);
44     double endTime = omp_get_wtime();
45     printf("%2.4f\n", endTime-startTime);
46
47     const int channel_nr = 1, stride_bytes = 0;
48     stbi_write_png("mandelbrot.png", X, Y, channel_nr, image, stride_bytes);
49     return EXIT_SUCCESS;
50 }


```

## 2.2 Measurement Method

The measurement was done on the LCC3 cluster by calling `sbatch job.sh mandelbrot 3`.

The same scripts as in exercise 1 are involved in the experiment.

## 2.3 Experiment Results



## 2.4 Discussion

The execution time drops significantly when the number of threads is increased. Considering that threads are completely independent here, this meets the expectations.