

Sheet 08

PS Parallel Programming

Patrick Wintner

May 12, 2025

1 Compiler Dependence Analysis

The dependence analysis of compilers is examined.

1.1 Source Code

```
1  #include <stdio.h>
2
3  #define SIZE 1024
4
5  int main(int argc, char** argv) {
6
7      int a[SIZE];
8      int b[SIZE];
9
10     for(int i = 0; i < SIZE; ++i) {
11         a[i] = argc;
12     }
13
14     for(int i = 0; i < SIZE; ++i) {
15         b[i] = a[i];
16     }
17
18     for(int i = 4; i < SIZE; ++i) {
19         a[i-4] = a[i];
20     }
```

```

21
22     for(int i = 1; i < SIZE-1; ++i) {
23         a[i] = a[i%argc];
24     }
25
26     // output data to prevent compiler from removing any code
27     for(int i = 0; i < SIZE; ++i) {
28         printf("%d ", a[i]);
29         printf("%d ", b[i]);
30     }
31     printf("\n");
32
33     return 0;
34 }

```

1.2 Makefile

```

CFLAGS = -O2 -ftree-vectorize -fopt-info-vec-all-internals

.PHONY: all
all: analysis

.PHONY: clean
clean:
    $(RM) analysis

analysis: analysis.c

```

1.3 Discussion of Compiler Output

The full compiler output can be found in the file 08/ex1/output.log.

Only the loop starting at line 10 has been successfully optimized. The last loop has not been optimized because the statement "clobbers memory" (which means here that the function call printf cannot be analyzed). The pretest loop starting at line 22 is analyzed using different vector modes and finally fails to vectorize the loop. The loop starting at line 10 has been optimized using 16 byte vectors. The loop starting at line 15 has not been vectorized because "more than one data ref in stmt" (probably because a and b could overlap). The loop starting at line 18 has also not been vectorized because the compiler previously recognized that the loop just performs a shift operation and has

replaced it with an appropriate builtin-function.

2 Investigation of Code Snippets

2.1 Safety of Parallelization

The safety of parallelization of the following code is examined.

```
1 void copy(double* x, double* y) {
2     for(int i = 0; i < 1024; i++) {
3         x[i] = y[i];
4     }
5 }
```

It is possible that the arrays `x` and `y` overlap each other. One way to parallelize this manually is by using a temporary array `z` with the `1024*sizeof(double)` and splitting the loop into two so that all read accesses from `y` (storing in `z`) are done in the first and all write accesses to `x` (reading from `z`) are done in the second. Both loops can run in parallel.

The function cannot be parallelized by the compiler, because the compiler needs the `__restrict__` keyword for both arguments, which implies that two pointers cannot point to overlapping memory regions.

2.1.1 Manually parallelized Loop

```
1 void copy(double* x, double* y) {
2     double* z = malloc(1024*sizeof(*z));
3     if(!z) {
4         return;
5     }
6     #pragma omp parallel for
7     for(int i = 0; i < 1024; i++) {
8         z[i] = y[i];
9     }
10    #pragma omp parallel for
11    for(int i = 0; i < 1024; i++) {
12        x[i] = z[i];
13    }
```

```

14     free(z);
15 }

```

2.2 Loop Normalization

The following loop should be normalisized.

```

1  for (int i=4; i<=N; i+=9) {
2      for (int j=0; j<=N; j+=5) {
3          A[i] = 0;
4      }
5  }

```

2.2.1 Normalisized Loop

```

1  for (int i=0; i<=(N-4)/9; i+=1) {
2      for (int j=0; j<=N/5; j+=1) {
3          A[i*9+4] = 0;
4      }
5  }

```

2.3 Parallelizability

It is examined if the following loop is parallelizable.

```

1  for(int i = 1; i < N; i++) {
2      for(int j = 1; j < M; j++) {
3          for(int k = 1; k < L; k++) {
4              a[i+1][j][k-1] = a[i][j][k] + 5;
5          }
6      }
7  }

```

The distance vector for all dependencies is $(1, 0, -1)$ and therefore the corresponding direction vector is $(<, =, >)$. Thus the outmostloop is not parallelizable. The dependency of the first inner loop is loop-independent, therefore that loop can be parallelized.

2.3.1 Parallelized Loop

```
1  for(int i = 1; i < N; i++) {  
2      #pragma omp parallel for  
3      for(int j = 1; j < M; j++) {  
4          for(int k = 1; k < L; k++) {  
5              a[i+1][j][k-1] = a[i][j][k] + 5;  
6          }  
7      }  
8  }
```